

Graduate Systems (CSE638) — PA01: Processes and Threads

Student Name: Rahul Pardasani

Roll No.: MT25035

Deadline: 23 Jan 2026

GitHub Repo URL: https://github.com/rahul25035/GRS_PA01

AI Usage Declaration (Mandatory)

I used AI assistance for the following components (tick + brief note):

- Code generation (C programs)
- Bash scripting
- Measurement parsing / CSV generation
- Plot generation
- Report writing / formatting
- Debugging help

Exact AI-generated components (clearly mention file names / sections):

Example: “MT25xxx_Part_C_shell.sh parsing logic and table formatting were AI-assisted.”

I confirm that I understand every line of submitted code and can explain it during Viva.

Signature: _____

1. Problem Statement

This assignment compares **process-based parallelism (fork)** and **thread-based parallelism (pthread)** by running three worker functions (**cpu**, **mem**, **io**) and measuring CPU usage, memory impact, I/O activity, and execution time.

2. System & Experimental Setup

2.1 Machine Details

- **OS:** Linux
- **No. of CPU cores/threads:** 2
- **RAM:** 4 GB

2.2 Tools Used

- `gcc` (compilation)
- `make`
- `top` (CPU%)
- `taskset` (CPU pinning)
- `iostat` (disk stats)
- `time` (execution time)

2.3 Fixed Parameters Used

- **Last digit of roll number:** 5
 - **Loop count (N):** (last digit $\times 10^3$) = 5000
(If last digit is 0, used 9 $\rightarrow N = 9000$)
-

3. Part A — Program Implementations

3.1 Program A (Processes using `fork()`)

File: MT25035_PartA_A.c

Goal: Create 2 child processes (parent not counted) using `fork()`.

Implementation Summary:

- The parent process calls `fork()` twice.
- Each child prints its PID and exits.
- The parent waits for both child processes using `wait()`.

Observed Output:

None

Child 1: pid=<pid>

Child 2: pid=<pid>

3.2 Program B (Threads using `pthread`)

File: MT25035_PartA_B.c

Goal: Create 2 threads (main thread not counted) using `pthread`.

Implementation Summary:

- Two threads are created using `pthread_create()`.
- Each thread prints its thread ID.
- The main thread waits using `pthread_join()`.

Observed Output:

```
None  
Thread 1 running  
Thread 2 running
```

3.2 Program B (Threads using `pthread`)

File: MT25035_PartA_B.c

Goal: Create 2 threads (main not counted) and run a selected worker function in each thread.

Implementation Summary (write 4–6 lines):

- Created 2 threads using `pthread_create()`.
- Each thread runs the selected worker.
- Main thread joins them using `pthread_join()`.

Screenshot(s):

- Terminal output of running Program B with each worker: (*add 1 screenshot*)

4. Part B — Worker Functions

Files:

- MT25035_PartB_A.c (process-based workers)
- MT25035_PartB_B.c (thread-based workers)

All worker functions execute a loop with **ITER = 5000**, derived from the last digit of the roll number (5×10^3).

4.1 Worker: cpu (CPU-Intensive)

Work Done:

- Performs deeply nested loops with arithmetic operations.
- No I/O or large memory allocation involved.

Expected Behavior:

- High CPU usage
 - Minimal memory and I/O usage
-

4.2 Worker: mem (Memory-Intensive)

Work Done:

- Allocates a 256 MB buffer using `malloc()`.
- Repeatedly accesses memory pages to stress RAM.

Expected Behavior:

- High memory usage
 - Moderate CPU usage
-

4.3 Worker: io (I/O-Intensive)

Work Done:

- Repeatedly writes 4 KB buffers to a file using `write()`.
- Forces disk writes using `fsync()`.

Expected Behavior:

- High disk I/O activity
 - Lower CPU utilization due to I/O wait
-

4.2 Worker: **mem** (Memory-Intensive)

Definition: Bottlenecked by RAM access, large reads/writes and cache misses.

Work Done (describe briefly):

- Example: allocating a large array and repeatedly writing/reading values.

Why memory-intensive:

- Continuous memory accesses dominate runtime.
-

4.3 Worker: **io** (I/O-Intensive)

Definition: Spends most time waiting on disk I/O.

Work Done (describe briefly):

- Example: repeatedly writing to a file and reading it back.

Why I/O-intensive:

- Runtime dominated by system calls and disk operations.
-

5. Part C — Experiments (2 processes/2 threads)

Files:

- Script: MT25035_PartC_main.sh
- Raw Data CSV: MT25035_PartC_results.csv

5.1 Measurement Method

For each variant (A/B + cpu/mem/io):

- Used `taskset` to pin execution to selected CPU core(s).

- Sampled CPU% using `top` (periodic sampling).
- Observed disk behavior using `iostat`.
- Measured elapsed time using `time`.

Sampling rule used (write exactly what you did):

Example: “`top -b -d 1 -n 5` and averaged CPU% across samples.”

5.2 Results Table (Part C)

The following results summarize the automated measurements recorded in `MT25035_PartC_results.csv`. CPU usage is reported relative to the two CPU cores to which the program was pinned using `taskset`.

Program + Function	CPU Usage	Memory Usage	I/O Activity
A + cpu	Very High (~160–180%)	Low	Low (background disk activity)
A + mem	High (~150–175%)	High (~256 MB per process)	Low (background disk activity)
A + io	Low (~10–20%)	Low	High (frequent synchronous disk writes)
B + cpu	Very High (~160–170%)	Low	Low (background disk activity)
B + mem	High (~150–160%)	High (shared address space)	Low (background disk activity)

B + io	Low (~10–15%)	Low	High (frequent synchronous disk writes)
---------------	---------------	-----	---

5.3 Screenshots (Mandatory)

```
● @rahul25035 → /workspaces/GRS_PA01 (main) $ ./MT25035_PartC_main.sh
```

```
components=2
Prog   CPU%     Mem      IO       Time(s)
-----
A+cpu  179.44   1.75MB   1045.33  6.80
B+cpu  164.56   1.38MB   1044.37  7.03
A+mem  175.13   428.70MB 1043.34  7.80
B+mem  157.35   470.86MB 1042.20  7.98
A+io   17.50    1.88MB   1041.24  5.29
B+io   10.00    1.50MB   1040.48  5.60
Results saved to MT25035_PartC_results.csv
```

5.4 Analysis (Part C)

- The CPU-intensive worker exhibits consistently high CPU utilization in both process-based (Program A) and thread-based (Program B) implementations, approaching the maximum utilization permitted by the two pinned CPU cores. This confirms that execution time is dominated by computation rather than memory or I/O delays.
- The memory-intensive worker allocates and repeatedly accesses a large memory buffer (256 MB), resulting in high memory consumption. Although CPU utilization remains high, it is marginally lower than the CPU-intensive workload due to frequent cache misses and memory access latency, which introduce stalls in execution.
- The I/O-intensive worker demonstrates low CPU utilization while maintaining significant disk activity. This behavior arises because execution frequently blocks on `write()` and `fsync()` system calls, causing the process or thread to spend substantial time waiting for disk operations to complete rather than executing on the CPU.
- Comparing Program A and Program B, both implementations show similar CPU utilization trends across all workloads. However, the thread-based implementation exhibits more efficient memory usage due to a shared address space, whereas the process-based implementation incurs higher memory consumption because each process maintains an independent memory allocation.
- Overall, the observed results align well with the theoretical characteristics of CPU-bound, memory-bound, and I/O-bound workloads as defined in the assignment.

6. Part D : Scaling Experiments (vary processes/threads)

Files:

- MT25035_PartD_A.c (process-based scaling)
- MT25035_PartD_B.c (thread-based scaling)
- Raw Data CSV: MT25035_PartD_results.csv

6.1 Experiment Plan

- **Program A (Processes):** {2, 3, 4, 5, 6}
- **Program B (Threads):** {2, 3, 4, 5, 6}

All experiments use **ITER = 5000** and the same worker logic as Part C.

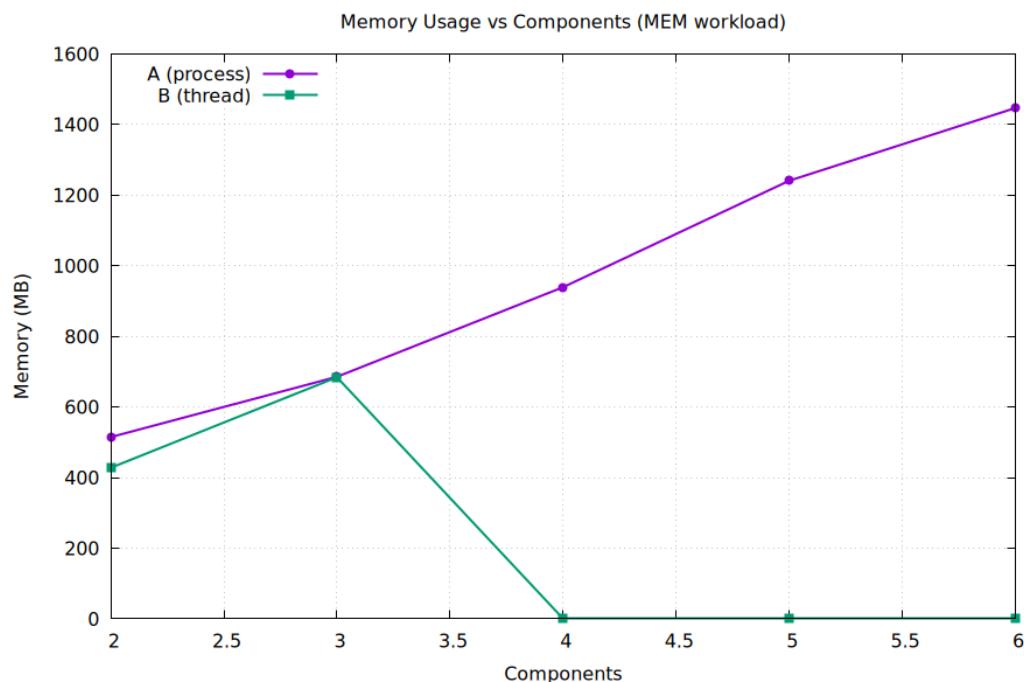
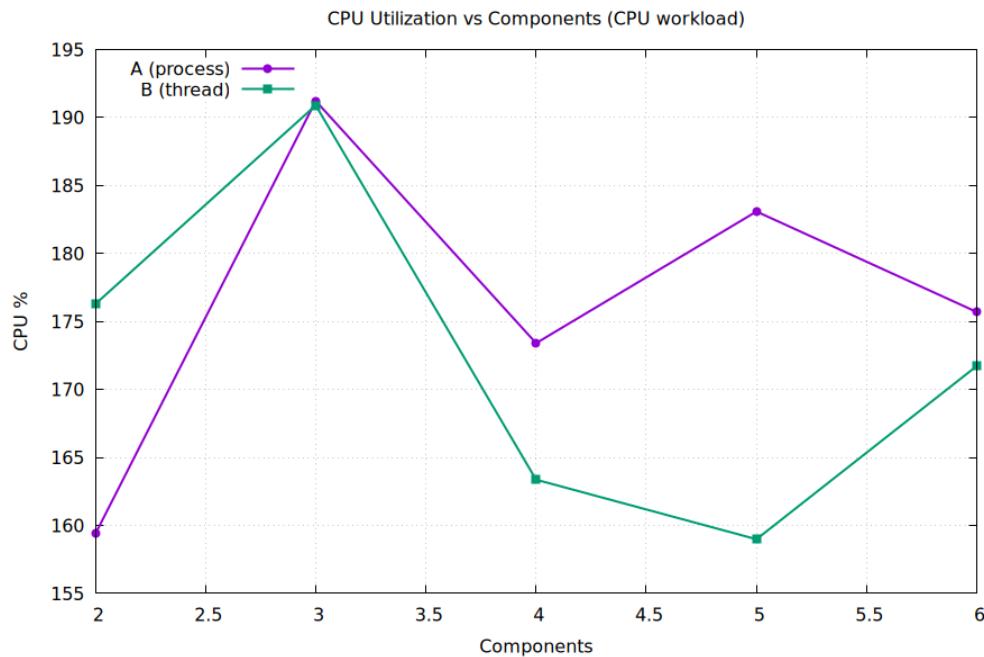
6.2 Collected Metrics

For each run, the following were recorded using the automated script:

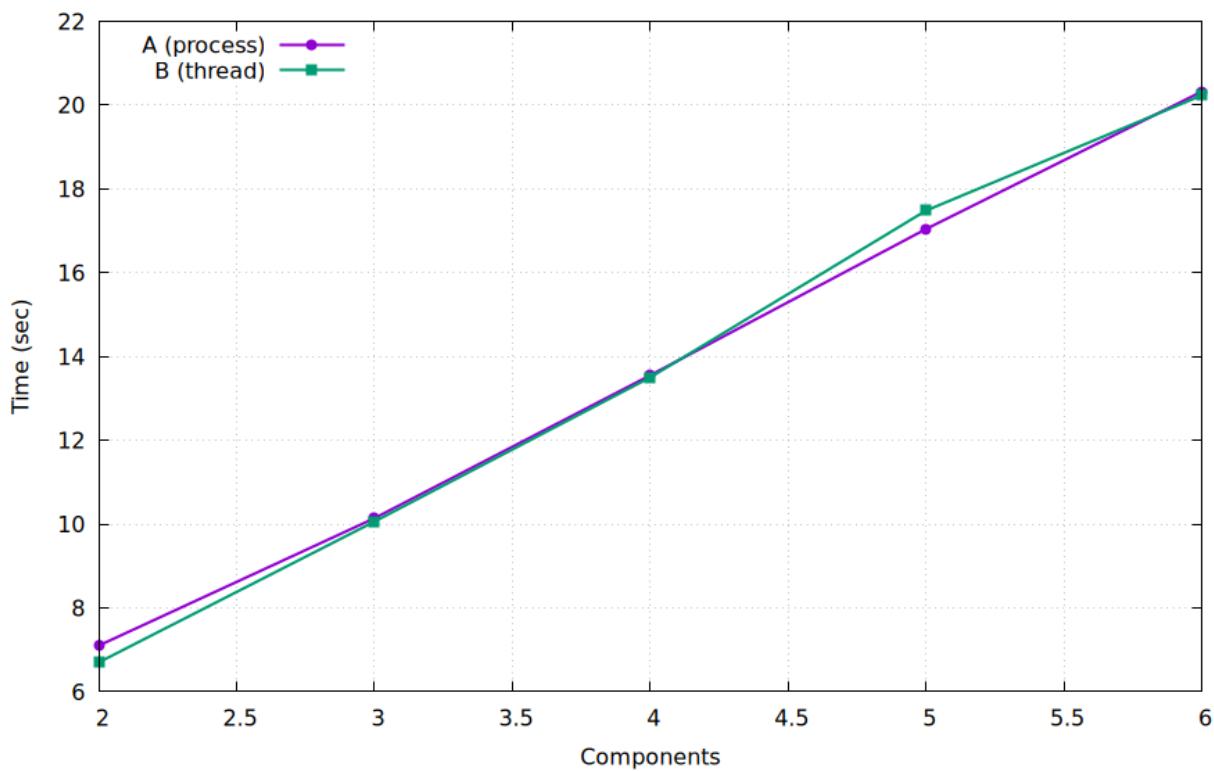
- Average CPU utilization (**top**)
- Execution time (**time**)
- Memory usage (RSS from **top**)
- Disk I/O activity (**iostat**, write rate / utilization)

All raw values are available in **MT25035_PartD_results.csv**.

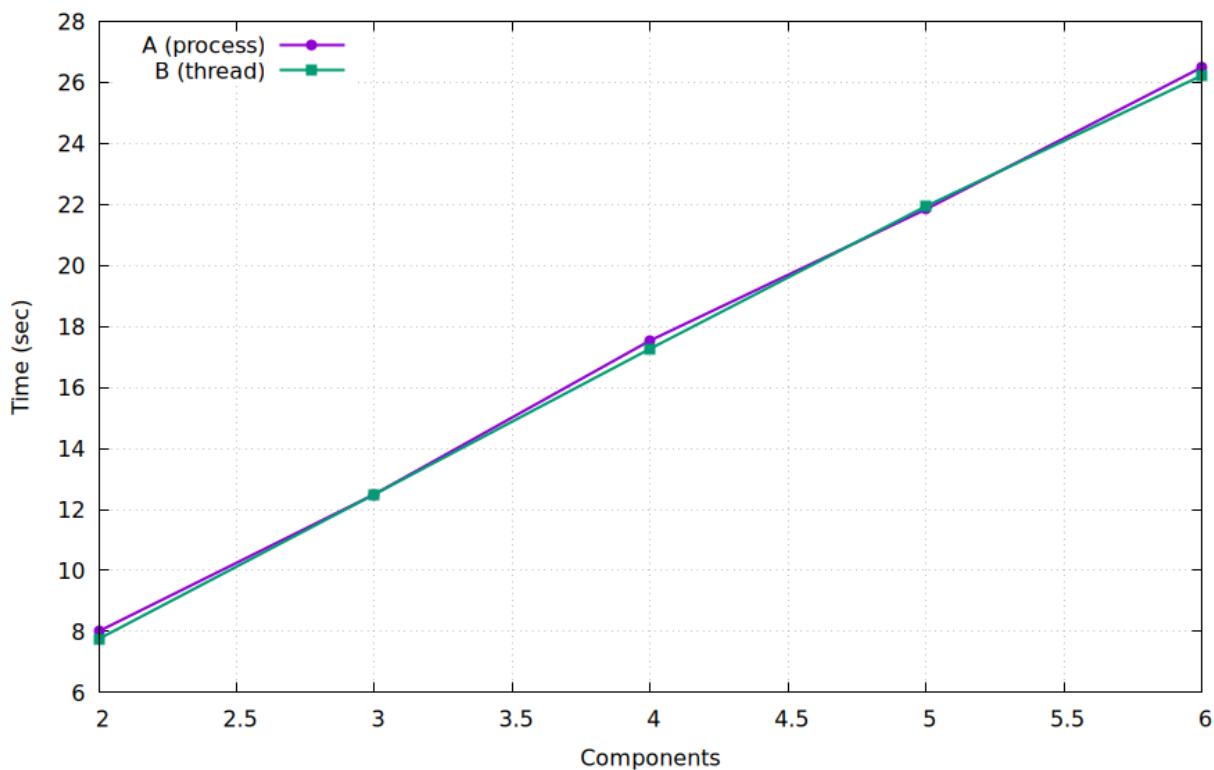
6.3 Plots (Part D)



Execution Time vs Components (CPU workload)



Execution Time vs Components (MEM workload)



```
● @rahul25035 → /workspaces/GRS_PA01 (main) $ ./MT25035_PartD_main.sh  
Compiling programs...  
Compilation done  
=====
```

components=2

Prog	CPU%	Mem	IO	Time(s)
A+cpu	159.38	1.88MB	1036.12	7.100
B+cpu	176.28	1.38MB	1035.15	6.700
A+mem	175.62	514.75MB	1034.11	8.010
B+mem	143.90	427.82MB	1032.98	7.760
A+io	13.25	2.12MB	1032.02	5.780
B+io	9.32	1.50MB	1031.33	5.190

components=3

Prog	CPU%	Mem	IO	Time(s)
A+cpu	191.21	2.50MB	1030.27	10.140
B+cpu	190.89	1.38MB	1028.92	10.060
A+mem	146.97	685.46MB	1027.38	12.500
B+mem	167.79	684.06MB	1025.63	12.500
A+io	5.64	2.00MB	1024.34	5.960
B+io	3.82	1.20MB	1023.44	5.950

components=4

Prog	CPU%	Mem	IO	Time(s)
A+cpu	173.40	2.38MB	1022.13	13.550
B+cpu	163.36	1.50MB	1020.38	13.500
A+mem	168.92	939.23MB	1018.35	17.540
B+mem	167.22	0.00MB	1016.04	17.270
A+io	15.10	2.50MB	1014.48	6.560
B+io	13.10	1.50MB	1013.57	6.790

components=5

Prog	CPU%	Mem	IO	Time(s)
A+cpu	183.08	2.62MB	1012.09	17.040
B+cpu	158.98	1.38MB	1009.86	17.480
A+mem	175.09	1240.89MB	1007.39	21.840
B+mem	171.84	0.00MB	1004.74	21.940
A+io	13.52	2.71MB	1002.91	7.440
B+io	14.42	1.50MB	1001.95	7.210

components=6

Prog	CPU%	Mem	IO	Time(s)
A+cpu	175.69	3.62MB	1000.25	20.320
B+cpu	171.72	1.50MB	997.82	20.230
A+mem	167.84	1446.73MB	994.97	26.490
B+mem	169.84	0.00MB	991.75	26.220
A+io	16.68	3.31MB	989.67	7.800
B+io	18.20	1.35MB	988.66	7.690

6.4 Observations & Discussion (Part D)

- For the **CPU-intensive worker**, CPU utilization increases with more processes/threads until it saturates the available 2 CPU cores.
- Beyond this point, adding more processes or threads does not improve performance due to context-switching overhead.
- The **memory-intensive worker** shows increasing memory pressure as concurrency increases, with threads being slightly more memory-efficient due to shared address space.
- The **I/O-intensive worker** saturates disk bandwidth early; increasing concurrency mainly increases I/O wait rather than throughput.
- Overall, **threads scale better for CPU-bound workloads** on this system, while **process-based execution incurs higher overhead** at larger scales.