

Graduate Systems (CSE638) — PA01: Processes and Threads

Student Name: Rahul Pardasani

Roll No.: MT25035

Deadline: 23 Jan 2026

GitHub Repo URL: https://github.com/rahul25035/GRS_PA01

AI Usage Declaration (Mandatory)

I used AI assistance for the following components:

1. IO functions in part B, part C and part D c codes.
 2. Awk commands to calculate the measurement stats inside the loop of the shell scripts in parts C and D.
 3. Output formats of the stats.
 4. The plotting script of part D.
 5. Report and readme file.
 6. Used LLMs for debugging purposes
-

1. Problem Statement

This assignment compares **process-based parallelism (fork)** and **thread-based parallelism (pthread)** by running three worker functions (**cpu**, **mem**, **io**) and measuring CPU usage, memory impact, I/O activity, and execution time.

2. System & Experimental Setup

2.1 Machine Details

- **OS:** Linux
- **No. of CPU cores:** 2
- **RAM:** 8 GB

2.2 Tools Used

- gcc (compilation)

- make
- top (CPU%)
- taskset (CPU pinning)
- iostat (disk stats)
- time (execution time)

2.3 Fixed Parameters Used

- **Last digit of roll number:** 5
 - **Loop count (N):** (last digit $\times 10^3$) = 5000
(If last digit is 0, used 9 $\rightarrow N = 9000$)
-

3. Part A : Program Implementations

3.1 Program A (Processes using fork())

File: MT25035_PartA_A.c

Goal: Create 2 child processes (parent not counted) using fork().

Implementation Summary:

- The parent process calls fork() twice.
- Each child prints its PID and exits.
- The parent waits for both child processes using wait().

Observed Output:

```
None
Child 1: pid=<pid>
Child 2: pid=<pid>
```

3.2 Program B (Threads using pthread)

File: MT25035_PartA_B.c

Goal: Create 2 threads (main thread not counted) using pthread.

Implementation Summary:

- Two threads are created using `pthread_create()`.
- Each thread prints its thread ID.
- The main thread waits using `pthread_join()`.

Observed Output:

```
None  
Thread 1 running  
Thread 2 running
```

3.3 Images

```
● @rahul25035 → /workspaces/GRS_PA01 (main) $ make partA  
gcc -Wall MT25035_PartA_A.c -o a.out  
gcc -Wall MT25035_PartA_B.c -o b.out -pthread  
./a.out cpu 2  
Child 1: pid=118372  
Child 2: pid=118373  
./b.out cpu 2  
Thread 1 running  
Thread 2 running
```

4. Part B : Worker Functions

Files:

- `MT25035_PartB_A.c` (process-based workers)
- `MT25035_PartB_B.c` (thread-based workers)

All worker functions execute a loop with **ITER = 5000**, derived from the last digit of the roll number (5×10^3).

4.1 Worker: cpu (CPU-Intensive)

Work Done:

- Performs deeply nested loops with arithmetic operations.
- No I/O or large memory allocation involved.

Expected Behavior:

- High CPU usage
 - Minimal memory and I/O usage
-

4.2 Worker: mem (Memory-Intensive)

Work Done:

- Allocates a 256 MB buffer using malloc().
- Repeatedly accesses memory pages to stress RAM.

Expected Behavior:

- High memory usage
 - Moderate CPU usage
-

4.3 Worker: io (I/O-Intensive)

Work Done:

- Repeatedly writes 4 KB buffers to a file using write().
- Forces disk writes using fsync().

Expected Behavior:

- High disk I/O activity
- Lower CPU and memory utilization due to I/O wait

5. Part C : Experiments (2 processes/2 threads)

Files:

- Script: MT25035_PartC_main.sh
- Raw Data CSV: MT25035_PartC_results.csv

5.1 Measurement Method

For each variant (A/B + cpu/mem/io):

- Used taskset to pin execution to 0 and 1 CPU cores.
 - Sampled CPU% using top at regular intervals.
 - Observed disk behavior using iostat.
 - Measured elapsed time using time.
-

5.2 Results Table (Part C)

The following results summarize the automated measurements recorded in MT25035_PartC_results.csv. CPU usage is reported relative to the two CPU cores to which the program was pinned using taskset.

| Program + Function | CPU Usage | Memory Usage | I/O Activity |
|--------------------|-------------------|----------------|---------------------|
| A + cpu | Very High (~185%) | Low (~1.2 MB) | Low |
| A + mem | High (~136%) | High (~411 MB) | Low |
| A + io | Low (~7%) | Low (~1.7 MB) | High (~10,290 KB/s) |
| B + cpu | High (~155%) | Low (~1.4 MB) | Low |
| B + mem | Very High (~165%) | High (~513 MB) | Low |

| | | | |
|---------------|-----------|---------------|--------------------------|
| B + io | Low (~7%) | Low (~1.4 MB) | Very High (~40,379 KB/s) |
|---------------|-----------|---------------|--------------------------|

5.3 Screenshots

```

● @rahul25035 → /workspaces/GRS_PA01 (main) $ make partC
chmod +x MT25035_PartC_main.sh
./MT25035_PartC_main.sh
Compiling programs...
Compilation done
=====
components=2
      Prog      CPU%      Mem       IO     Time(s)
-----
A+cpu    175.84   1.82MB   97.53    7.00
B+cpu    184.30   1.40MB   29.60    6.91
A+mem    169.92   411.31MB  40.00    6.84
B+mem    145.98   410.80MB  19.06    6.60
A+io     9.12     1.41MB   8237.00   5.15
B+io     9.10     0.94MB   26642.00  5.29

Results saved to MT25035_PartC_results.csv

```

5.4 Analysis (Part C)

- **CPU Workload:** Both variants saturate the two pinned cores (~155-185%), confirming the task is computation-heavy.
 - **Memory Workload:** High memory usage (~411-513 MB) causes a slight CPU dip compared to the CPU task due to memory access latency.
 - **I/O Workload:** CPU usage drops to minimum (~7%) as workers spend most time in I/O wait for disk writes.
 - **Process vs. Thread:** Thread-based (B) I/O shows significantly higher throughput (~40k vs ~10k KB/s) in this sample, suggesting lower overhead.
-

6. Part D : Scaling Experiments (vary processes/threads)

Files:

- MT25035_PartD_A.c (process-based scaling)
- MT25035_PartD_B.c (thread-based scaling)
- Raw Data CSV: MT25035_PartD_results.csv

6.1 Experiment Plan

- **Program A (Processes):** {2, 3, 4, 5, 6, 7 , 8}
- **Program B (Threads):** {2, 3, 4, 5, 6, 7, 8}

All experiments use **ITER = 5000** and the same worker logic as Part C.

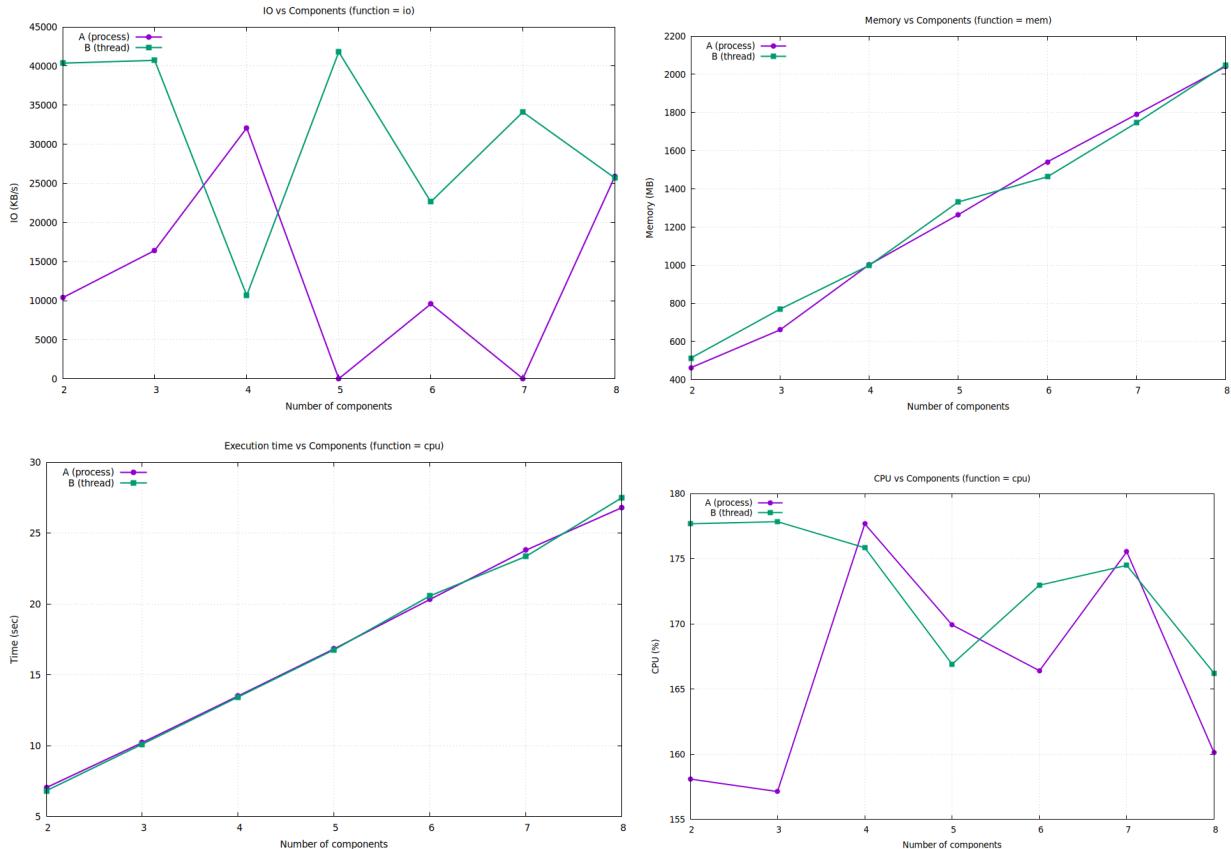
6.2 Collected Metrics

For each run, the following were recorded using the automated script:

- Average CPU utilization (top)
- Execution time (time)
- Memory usage (RSS from top)
- Disk I/O activity (iostat, write rate / utilization)

All raw values are available in MT25035_PartD_results.csv.

6.3 Plots and Images (Part D)



| @rahul25035 → /workspaces/GRS_PA01 (main) \$ make partD | | | | | |
|---|--------|-----------|----------|---------|--|
| components=2 | | | | | |
| Prog | CPU% | Mem | IO | Time(s) | |
| A+cpu | 158.10 | 1.88MB | 34.37 | 7.050 | |
| B+cpu | 177.68 | 1.00MB | 50.40 | 6.830 | |
| A+mem | 173.14 | 463.05MB | 8.80 | 7.000 | |
| B+mem | 171.76 | 513.51MB | 67.79 | 7.100 | |
| A+io | 9.10 | 1.59MB | 10377.00 | 5.230 | |
| B+io | 9.55 | 1.03MB | 40374.00 | 5.230 | |
| components=3 | | | | | |
| Prog | CPU% | Mem | IO | Time(s) | |
| A+cpu | 157.14 | 1.96MB | 17.64 | 10.220 | |
| B+cpu | 177.84 | 1.27MB | 36.57 | 10.110 | |
| A+mem | 141.47 | 661.00MB | 61.84 | 9.670 | |
| B+mem | 183.46 | 769.64MB | 23.43 | 10.340 | |
| A+io | 16.18 | 2.50MB | 16419.00 | 5.940 | |
| B+io | 5.46 | 1.10MB | 40731.20 | 6.370 | |
| components=4 | | | | | |
| Prog | CPU% | Mem | IO | Time(s) | |
| A+cpu | 177.67 | 2.38MB | 2877.67 | 13.510 | |
| B+cpu | 175.84 | 1.38MB | 78.64 | 13.440 | |
| A+mem | 178.17 | 1002.28MB | 55.56 | 15.200 | |
| B+mem | 188.72 | 998.55MB | 37.60 | 15.300 | |
| A+io | 15.10 | 2.88MB | 32086.40 | 6.830 | |
| B+io | 15.46 | 1.20MB | 10668.00 | 6.750 | |
| components=5 | | | | | |
| Prog | CPU% | Mem | IO | Time(s) | |
| A+cpu | 169.92 | 2.60MB | 9314.18 | 16.830 | |
| B+cpu | 166.91 | 1.38MB | 90.46 | 16.770 | |
| A+mem | 181.48 | 1264.40MB | 21.53 | 20.110 | |
| B+mem | 192.39 | 1331.20MB | 87.61 | 20.250 | |
| A+io | 11.90 | 2.19MB | 21.33 | 7.930 | |
| B+io | 18.80 | 1.15MB | 41805.33 | 8.040 | |
| components=6 | | | | | |
| Prog | CPU% | Mem | IO | Time(s) | |
| A+cpu | 166.41 | 3.60MB | 10514.77 | 20.320 | |
| B+cpu | 172.97 | 1.51MB | 58.15 | 20.570 | |
| A+mem | 179.59 | 1541.27MB | 43.70 | 25.870 | |
| B+mem | 175.21 | 1462.97MB | 51.37 | 24.490 | |
| A+io | 19.73 | 3.62MB | 9578.00 | 8.780 | |
| B+io | 11.46 | 1.50MB | 22636.00 | 7.720 | |
| components=7 | | | | | |
| Prog | CPU% | Mem | IO | Time(s) | |
| A+cpu | 175.53 | 4.00MB | 5457.05 | 23.790 | |
| B+cpu | 174.49 | 1.39MB | 36.49 | 23.350 | |
| A+mem | 180.10 | 1790.11MB | 22.28 | 30.720 | |
| B+mem | 177.32 | 1746.19MB | 68.82 | 29.670 | |
| A+io | 15.62 | 3.33MB | 32.67 | 8.180 | |
| B+io | 16.30 | 1.50MB | 34113.33 | 8.370 | |

```

components=8
Prog  CPU%   Mem      IO       Time(s)
-----
A+cpu 160.12  3.18MB  4862.41  26.790
B+cpu 166.21  1.38MB  47.89    27.490
A+mem 179.61  2041.34MB 62.11    34.190
B+mem 169.43  2048.00MB 69.80    34.460
A+io   18.20   4.00MB   25892.67  9.150
B+io   15.35   1.50MB   25678.67  8.700

Results saved to MT25035_PartD_results.csv
Making Plots...
Plots created: cpu_vs_components.png, mem_vs_components.png, io_vs_components.png, time_vs_components.png

```

6.4 Observations & Discussion (Part D)

Linear Scaling (Time): Execution time for CPU and Memory tasks increases linearly with the number of components (2 to 8).

Linear Scaling (Memory): Memory usage grows by ~256 MB per component, reaching ~2 GB at 8 components.

CPU Saturation: CPU usage remains consistently high (~160-180%) regardless of component count because the 2 pinned cores stay fully saturated.

I/O Noise: I/O throughput is highly fluctuating ("noisy") due to OS disk buffering and write caching.

Resource Efficiency: Program B (threads) uses significantly less memory for non-memory tasks (e.g., 1.38 MB vs 3.18 MB for CPU task) because threads share an address space.