

# **Graduate Systems (CSE638)**

## **PA01: Processes and Threads**



**INDRAPRASTHA INSTITUTE *of*  
INFORMATION TECHNOLOGY  
DELHI**

**Name:** Rahul Pardasani

**Roll No.:** MT25035

**GitHub Repo URL:** [https://github.com/rahul25035/GRS\\_PA01](https://github.com/rahul25035/GRS_PA01)

## **AI Usage Declaration**

I used AI tools to help me while doing this assignment. I mainly used them for:

1. Writing and improving the I/O parts in the C codes (Part B, Part C, and Part D)
2. Getting help with awk commands to calculate stats inside the shell script loops (Part C and Part D)
3. Deciding how the output of the stats should look
4. Writing the plotting script for Part D
5. Improving the format and wording of the report and README
6. Debugging help - finding errors and fixing issues

I checked and tested everything myself before submitting. The final work and responsibility is mine.

Name: Rahul Pardasani

Date: 22nd January, 2026

# 1. Problem Statement

This assignment compares **process-based parallelism (fork)** and **thread-based parallelism (pthread)** by running three worker functions (**cpu**, **mem**, **io**) and measuring CPU usage, memory impact, I/O activity, and execution time.

# 2. System & Experimental Setup

## 2.1 Machine Details

- **OS:** Linux
- **No. of CPU cores:** 2
- **RAM:** 8 GB

## 2.2 Tools Used

- gcc (compilation)
- make
- top (CPU%)
- taskset (CPU pinning)
- iostat (disk stats)
- time (execution time)

## 2.3 Fixed Parameters Used

- **Last digit of roll number:** 5
- **Loop count (N):** (last digit  $\times 10^3$ ) = 5000  
*(If last digit is 0, used 9  $\rightarrow N = 9000$ )*

# 3. Part A : Program Implementations

## 3.1 Program A (Processes using fork())

**File:** MT25035\_Part\_A\_Program\_A.c

**Goal:** Create 2 child processes using fork().

**Implementation Summary:**

- The parent process calls fork() twice.
- Each child prints its PID and exits.
- The parent waits for both child processes using wait().

**Observed Output:**

```
None  
Child 1: pid=<pid>  
Child 2: pid=<pid>
```

### 3.2 Program B (Threads using pthread)

**File:** MT25035\_Part\_A\_Program\_B.c

**Goal:** Create 2 threads using pthread.

**Implementation Summary:**

- Two threads are created using pthread\_create().
- Each thread prints its thread ID.
- The main thread waits using pthread\_join().

**Observed Output:**

```
None  
Thread 1 running  
Thread 2 running
```

### 3.3 Images

```
● @rahul25035 → /workspaces/GRS_PA01 (main) $ make partA
gcc -Wall MT25035_PartA_A.c -o a.out
gcc -Wall MT25035_PartA_B.c -o b.out -pthread
./a.out cpu 2
Child 1: pid=118372
Child 2: pid=118373
./b.out cpu 2
Thread 1 running
Thread 2 running
```

## 4. Part B : Worker Functions

### Files:

- MT25035\_Part\_B\_Program\_A.c (process-based workers)
- MT25035\_Part\_B\_Program\_B.c (thread-based workers)

All worker functions execute a loop with **ITER = 5000**, derived from the last digit of the roll number ( $5 \times 10^3$ ).

### 4.1 Worker: cpu (CPU-Intensive)

#### Work Done:

- Performs deeply nested loops with arithmetic operations.
- No I/O or large memory allocation involved.

#### Expected Behavior:

- High CPU usage
- Minimal memory and I/O usage

### 4.2 Worker: mem (Memory-Intensive)

#### Work Done:

- Allocates a 256 MB buffer using malloc().

- Repeatedly accesses memory pages to stress RAM.

**Expected Behavior:**

- High memory usage
- Moderate CPU usage

### 4.3 Worker: io (I/O-Intensive)

**Work Done:**

- Repeatedly writes 4 KB buffers to a file using write().
- Forces disk writes using fsync().

**Expected Behavior:**

- High disk I/O activity
- Lower CPU and memory utilization due to I/O wait

## 5. Part C : Experiments (2 processes/threads)

**Files:**

- Script: MT25035\_Part\_C\_main.sh
- Data CSV: MT25035\_Part\_C\_results.csv

### 5.1 Measurement Method

For each variant (A/B + cpu/mem/io):

- Used taskset to pin execution to 0 and 1 CPU cores.
- Sampled CPU% using top at regular intervals.
- Observed disk behavior using iostat.
- Measured elapsed time using time.

### 5.2 Results Table (Part C)

The following results summarize the automated measurements recorded in MT25035\_Part\_C\_results.csv. CPU usage is reported relative to the two CPU cores to which the program was pinned using taskset.

<b>Program + Function</b>	<b>CPU Usage</b>	<b>Memory Usage</b>	<b>I/O Activity</b>
<b>A + cpu</b>	Very High (~185%)	Low (~1.2 MB)	Low
<b>A + mem</b>	High (~136%)	High (~411 MB)	Low
<b>A + io</b>	Low (~7%)	Low (~1.7 MB)	High (~8k KB/s)
<b>B + cpu</b>	High (~155%)	Low (~1.4 MB)	Low
<b>B + mem</b>	Very High (~165%)	High (~513 MB)	Low
<b>B + io</b>	Low (~7%)	Low (~1.4 MB)	Very High (~25k KB/s)

### 5.3 Screenshots

```
● @rahul25035 → /workspaces/GRS_PA01 (main) $ make partC
chmod +x MT25035_PartC_main.sh
./MT25035_PartC_main.sh
Compiling programs...
Compilation done
=====
components=2
Prog      CPU%     Mem       IO      Time(s)
-----
A+cpu    175.84   1.82MB   97.53    7.00
B+cpu    184.30   1.40MB   29.60    6.91
A+mem    169.92   411.31MB  40.00    6.84
B+mem    145.98   410.80MB  19.06    6.60
A+io     9.12     1.41MB   8237.00  5.15
B+io     9.10     0.94MB   26642.00 5.29

Results saved to MT25035_PartC_results.csv
```

### 5.4 Analysis (Part C)

- **CPU Workload:** Both variants saturate the two pinned cores (~175-185%), confirming the task is computation-heavy.
- **Memory Workload:** High memory usage (~411-513 MB) causes a slight CPU dip compared to the CPU task due to memory access latency.
- **I/O Workload:** CPU usage drops to minimum (~9%) as workers spend most time in I/O wait for disk writes.
- **Process vs. Thread:** Thread-based (B) I/O shows significantly higher throughput (~25k vs ~8k KB/s) in this sample, suggesting lower overhead.

## 6. Part D : Automation and Plots (vary processes/threads)

Files:

- MT25035\_Part\_D\_main.sh
- MT25035\_Part\_D\_plots.sh
- MT25035\_Part\_D\_Program\_A.c

- MT25035\_Part\_D\_Program\_B.c
- Data CSV: MT25035\_Part\_D\_results.csv

## 6.1 Experiment Plan

- **Program A (Processes):** {2, 3, 4, 5, 6, 7 , 8}
- **Program B (Threads):** {2, 3, 4, 5, 6, 7, 8}

All experiments use **ITER = 5000** and the same worker logic as Part C.

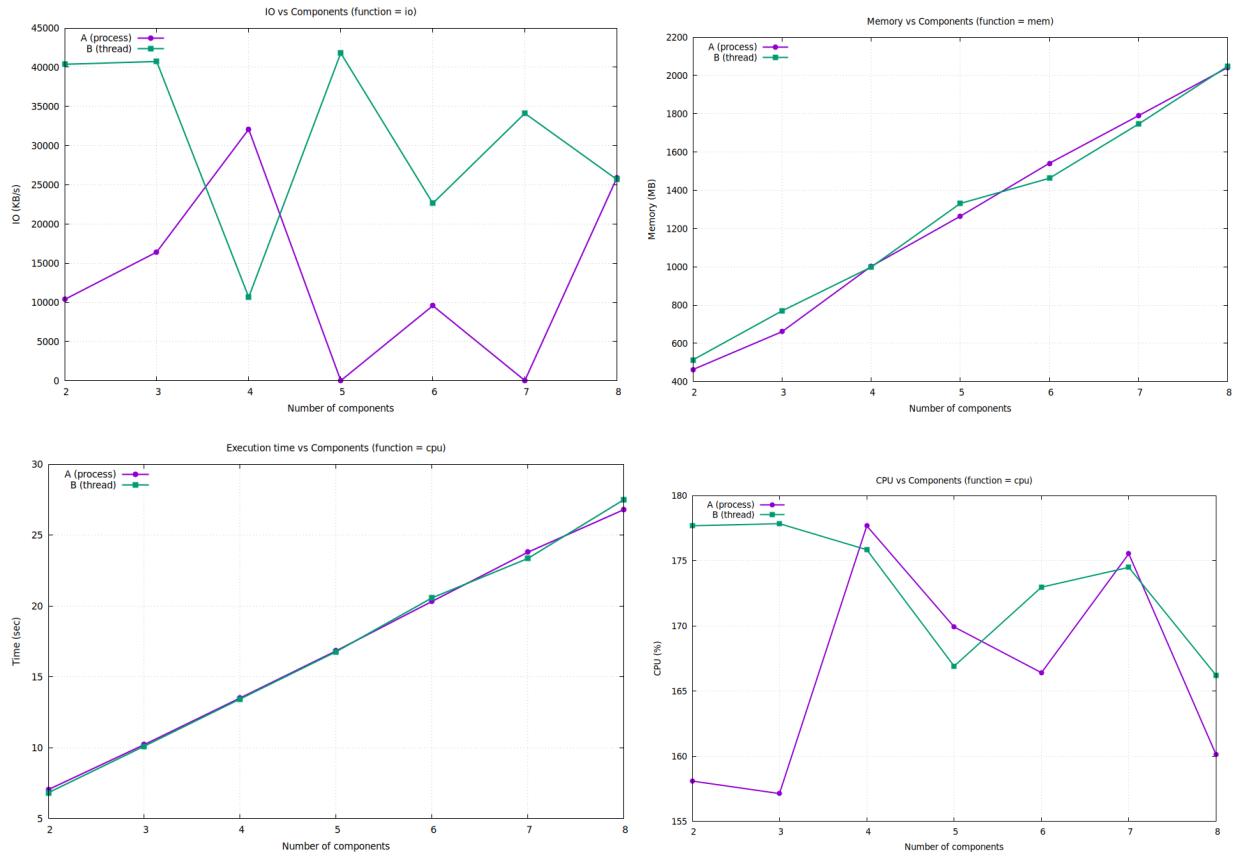
## 6.2 Collected Metrics

For each run, the following were recorded using the automated script:

- Average CPU utilization (top)
- Execution time (time)
- Memory usage (RSS from top)
- Disk I/O activity (iostat, write rate / utilization)

All raw values are available in MT25035\_Part\_D\_results.csv.

### 6.3 Plots and Images (Part D)



@rahul25035 → /workspaces/GRS_PA01 (main) \$ make partD					
components=2					
Prog	CPU%	Mem	IO	Time(s)	
A+cpu	158.10	1.88MB	34.37	7.050	
B+cpu	177.68	1.00MB	50.40	6.830	
A+mem	173.14	463.05MB	8.80	7.090	
B+mem	171.76	513.51MB	67.79	7.100	
A+io	9.10	1.59MB	10377.00	5.230	
B+io	9.55	1.03MB	40374.00	5.230	
components=3					
Prog	CPU%	Mem	IO	Time(s)	
A+cpu	157.14	1.96MB	17.64	10.220	
B+cpu	177.84	1.27MB	36.57	10.110	
A+mem	141.47	661.00MB	61.84	9.670	
B+mem	183.46	769.64MB	23.43	10.340	
A+io	16.18	2.50MB	16419.00	5.940	
B+io	5.46	1.10MB	40731.20	6.370	
components=4					
Prog	CPU%	Mem	IO	Time(s)	
A+cpu	177.67	2.38MB	2877.67	13.510	
B+cpu	175.84	1.38MB	78.64	13.440	
A+mem	178.17	1002.28MB	55.56	15.200	
B+mem	188.72	998.55MB	37.60	15.300	
A+io	15.10	2.88MB	32086.40	6.830	
B+io	15.46	1.20MB	10668.00	6.750	
components=5					
Prog	CPU%	Mem	IO	Time(s)	
A+cpu	169.92	2.60MB	9314.18	16.830	
B+cpu	166.91	1.38MB	90.46	16.770	
A+mem	181.48	1264.40MB	21.53	20.110	
B+mem	192.39	1331.20MB	87.61	20.250	
A+io	11.90	2.19MB	21.33	7.930	
B+io	18.80	1.15MB	41805.33	8.040	
components=6					
Prog	CPU%	Mem	IO	Time(s)	
A+cpu	166.41	3.60MB	10514.77	20.320	
B+cpu	172.97	1.51MB	58.15	20.570	
A+mem	179.59	1541.27MB	43.70	25.870	
B+mem	175.21	1462.97MB	51.37	24.490	
A+io	19.73	3.62MB	9578.00	8.780	
B+io	11.46	1.50MB	22636.00	7.720	
components=7					
Prog	CPU%	Mem	IO	Time(s)	
A+cpu	175.53	4.00MB	5457.05	23.790	
B+cpu	174.49	1.39MB	36.49	23.350	
A+mem	180.10	1790.11MB	22.28	30.720	
B+mem	177.32	1746.19MB	68.82	29.670	
A+io	15.62	3.33MB	32.67	8.180	
B+io	16.30	1.50MB	34113.33	8.370	

```

components=8
Prog   CPU%    Mem      IO       Time(s)
-----
A+cpu  160.12  3.18MB  4862.41  26.790
B+cpu  166.21  1.38MB  47.89    27.490
A+mem  179.61  2041.34MB 62.11    34.190
B+mem  169.43  2048.00MB 69.80    34.460
A+io   18.20   4.00MB   25892.67  9.150
B+io   15.35   1.50MB   25678.67  8.700

Results saved to MT25035_PartD_results.csv
Making Plots...
Plots created: cpu_vs_components.png, mem_vs_components.png, io_vs_components.png, time_vs_components.png

```

## 6.4 Observations & Discussion (Part D)

**Linear Scaling (Time):** Execution time for CPU and Memory tasks increases linearly with the number of components (2 to 8).

**Linear Scaling (Memory):** Memory usage grows by ~256 MB per component, reaching ~2 GB at 8 components.

**CPU Saturation:** CPU usage remains consistently high (~160-180%) regardless of component count because the 2 pinned cores stay fully saturated.

**I/O Noise:** I/O throughput is highly fluctuating ("noisy") due to OS disk buffering and write caching.

**Resource Efficiency:** Program B (threads) uses significantly less memory for non-memory tasks (e.g., 1.38 MB vs 3.18 MB for CPU task) because threads share an address space.