

# PA01 — Processes vs Threads

**Name:** Rahul Pardasani

**Course:** Graduate Systems (CSE638)

**Assignment:** PA01 — Processes and Threads

## 1. Objective

Compare process-based (A) and thread-based (B) implementations for three workloads (CPU-bound, Memory-bound, I/O-bound) by measuring CPU%, memory (MB), I/O activity and wall-clock time while varying the number of components (2–6). ITER = **5000**.

## 2. Implementation (brief)

- **A.c** — creates N child processes with fork(). Each child runs one worker: cpu\_func, mem\_func, or io\_func. Memory allocations are per-process.
- **B.c** — creates N pthreads. Each thread runs one worker in the same address space.
- **Workers:**
  - cpu — nested loops to burn CPU cycles.
  - mem — allocates ~256 MB and touches pages repeatedly.
  - io — writes 4 KB buffers and calls fsync() repeatedly.

## 3. Measurement method (brief)

main.sh automates experiments and logs results.csv. For each run: pinned to CPUs 0–1 (taskset), sampled per-second top for CPU and RSS, sampled iostat for I/O activity, and measured elapsed time with GNU /usr/bin/time. Data was aggregated into averages per run.

Note: the iostat field used is an I/O activity indicator (interpret as relative I/O load, not exact KB/s).

## 4. Key results (averages across component runs)

<b>Program</b>	<b>Function</b>	<b>Avg CPU%</b>	<b>Avg Mem (MB)</b>	<b>Avg Time (s)</b>
A	cpu	168.18	2.43	13.60
B	cpu	165.15	1.41	13.55
A	mem	160.02	1003.50	16.86
B	mem	162.80	246.75	17.07
A	io	13.59	2.55	6.72
B	io	11.53	1.47	6.62

## 5. Outcomes per combination (concise, clear)

### A + cpu (processes, CPU-bound)

- **What happened:** Very high CPU% ( $\approx 168\%$ ). Low memory usage ( $\sim 2\text{--}3$  MB RSS). Execution time grows as components increase (e.g.,  $\sim 6.9\text{s}$  at 2 components  $\rightarrow \sim 20.3\text{s}$  at 6) because only two cores are pinned, so more processes contend for CPU.
- **Conclusion:** Expected, CPU is the bottleneck; additional processes share the same cores, increasing contention.

### B + cpu (threads, CPU-bound)

- **What happened:** Similar to A: high CPU% ( $\sim 165\%$ ) and similar times ( $\sim 13.6\text{s}$  mean). Memory negligible.
- **Conclusion:** Threads and processes perform nearly the same on pure CPU-bound work when cores are saturated.

### A + mem (processes, Memory-bound)

- **What happened:** High CPU% ( $\sim 160\%$ ) and large RSS ( $\approx 1000$  MB average) because each process allocates its own  $\sim 256$  MB buffer and duplicates memory across processes.
- **Conclusion:** Expected, processes duplicate memory, so total RAM grows with component count. Time is driven by memory access and bandwidth.

## **B + mem (threads, Memory-bound)**

- **What happened:** CPU% remains high (~163%) but RSS is much lower (~247 MB average) because threads share one address space (memory not duplicated).
- **Conclusion:** Threads are significantly more memory-efficient for this workload while run times remain comparable.

## **A + io (processes, I/O-bound)**

- **What happened:** Low CPU% (~13.6%), low RSS, and higher I/O activity; elapsed times ≈6–8s and don't improve much with more components.
- **Conclusion:** Disk I/O is the bottleneck; processes spend most time waiting on I/O.

## **B + io (threads, I/O-bound)**

- **What happened:** Similar to A: low CPU% (~11.5%), low memory, similar elapsed times (~6.6s).
- **Conclusion:** I/O-bound behavior is dominated by the storage subsystem; neither processes nor threads gain much advantage.

## **6. Overall analysis (two-line summary)**


- **CPU-bound:** Both designs saturate the available cores and show similar performance.
- **Memory-bound:** Threads win on memory usage; run times are similar.
- **I/O-bound:** Both are limited by disk; little difference in runtime.

## **7. Important measurement caveat**


The I/O column in results.csv is an activity indicator derived from iostat parsing used in the script. Treat it as relative I/O load (higher → more I/O), not as exact KB/s unless you change the iostat fields extracted.



## **8. Plots**

CPU Utilization vs Components (CPU workload)



Memory Usage vs Components (MEM workload)





## 9. Files to submit (already prepared)

A.c, B.c, main.sh, plots.sh, Makefile, results.csv, and generated plots (time\_cpu.png, time\_mem.png, cpu\_usage.png, mem\_usage.png).

## **10. AI usage declaration**

I used ChatGPT to help write and polish this report. I can explain and defend every line of code, the measurement method, and the claims in the viva.